


A hand-drawn illustration featuring a network of interconnected nodes and lines. The nodes are represented by various colored shapes: a blue cube, a yellow house-like shape, a red cube, a green cube, a green rectangle, a yellow cube, and a blue cube. These shapes are connected by thin black lines, forming a complex web. In the center of the image, the word "Deadlock" is written in a large, bold, black sans-serif font. The background is a light, textured grey.




- Deadlock
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Dining philosopher's problem



Deadlock

- Permanent blocking of a set of processes that compete for system resources
- A set of processes is deadlocked when each of the process in the set is blocked awaiting an event that can only be triggered by another blocked process in the set.
- Involve conflicting needs for resources by two or more processes



Deadlock






Figure 6.1 Illustration of Deadlock




- Under normal mode of operation, a process may utilize a resource in following sequence:
 1. Request
 2. Use
 3. Release






System Table


- Records whether each resource is free or allocated
- If allocated then to which process
- If a process requests a resource currently allocated to another process then it is added to a queue of processes waiting for that resource.






Consumable Resources


- Created (produced) and destroyed (consumed)
- There is no limit on no of consumable resources of a particular type.
- Information in I/O buffers, Interrupts, signals, messages






Reusable Resources

- Used by only one process at a time and not depleted by that use
- Processes obtain resources that they later release for reuse by other processes
- Processors, I/O channels, main and secondary memory, devices, and data structures such as files, databases
- Deadlock occurs if each process holds one resource and requests the other





Resource Allocation Graphs


- Directed graph that depicts a state of the system of resources and processes



(a) Resource is requested




(b) Resource is held






Resource Allocation Graphs



(c) Circular wait



(d) No deadlock



Resource Allocation Graphs




Figure 7.1 Resource-allocation graph.



Resource Allocation Graphs






Figure 6.6 Resource Allocation Graph for Figure 6.1b






Conditions for Deadlock


- Mutual exclusion
 - Only one process may use a resource at a time
- Hold-and-wait
 - A process may hold allocated resources while awaiting assignment of others






Conditions for Deadlock

- No preemption
 - No resource can be forcibly removed from a process holding it
- Circular wait
 - A closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain






Possibility of Deadlock


- Mutual Exclusion
- No preemption
- Hold and wait






Existence of Deadlock

- Mutual Exclusion
- No preemption
- Hold and wait
- Circular wait





Handling Deadlocks

- Prevent
- Avoid
- Detect





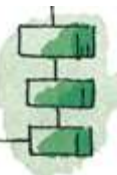
Deadlock Prevention


Indirect Method

Direct Method

Prevent occurrence of
conditions 1 to 3


Prevent occurrence of condition 4
i.e. Circular Wait






Deadlock Prevention


- Mutual Exclusion
 - Must be supported by the OS
 - Sharable resources
 - Non Sharable resources






Hold and Wait

- Require a process request all of its required resources at one time
- A process may request some resources and may use them.
- Before it can request any other resource it must release all the resources that is currently allocated.
- Disadvantages
 - Starvation
 - Low resource utilization
 - No prior knowledge of resources required





No Preemption


- Process must release resource and request again
- OS may preempt a process to require it releases its resources






Circular Wait

- Define a linear ordering of resource types
- Circular wait can be prevented by defining a linear ordering of resource types and to require that each process requires resource in an increasing order of enumeration.




Deadlock Avoidance




Requires that the system has some additional *a priori* information available

- Simplest and most useful model requires that each process declare the **maximum number** of resources of each type that it may need
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes
- Deadlock Avoidance Algorithms —
 - Safe State Algorithm
 - Banker's Algorithm




Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of ALL the processes in the systems such that for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
- That is:
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on




Basic Facts

- If a system is in safe state \Rightarrow no deadlocks
- If a system is in unsafe state \Rightarrow possibility of deadlock
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.



Safe, Unsafe, Deadlock State



Safe State


- Example : Consider 3 processes P1, P2, P3 and one resource R1 with 12 instances

Process	Max. Need	Current Allocation	Need	Available (After process execution)
P1	10	5	5	10
P2	4	2	2	5
P3	9	2	7	12
Total :	23	9		
Free resources :		3		

Free resources = Total resources-current allocation = $12 - 9 = 3$


Safe sequence = $\langle P2, P1, P3 \rangle$

Available = Allocated free resources + current allocation



Banker's Algorithm

- Multiple instances
- Each process must a priori claim maximum use
- When a process requests a resource it may have to wait
- When a process gets all its resources it must return them in a finite amount of time



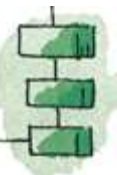


Data Structures for the Banker's Algorithm

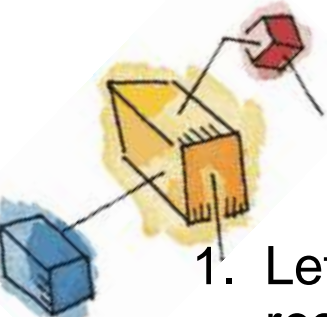
Let n = number of processes, and m = number of resources types.

- **Available:** Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available
- **Max:** $n \times m$ matrix. If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j
- **Allocation:** $n \times m$ matrix. If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j
- **Need:** $n \times m$ matrix. If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task

$$Need[i, j] = Max[i, j] - Allocation[i, j]$$



Safety Algorithm

- 
1. Let **Work** and **Finish** be vectors of length m and n , respectively. Initialize:

Work = Available

Finish [i] = **false** for $i = 0, 1, \dots, n-1$

2. Find an i such that both:


(a) **Finish** [i] = **false**

(b) **Need** _{i} ≤ **Work**

If no such i exists, go to step 4

3. **Work = Work + Allocation** _{i}
Finish [i] = **true**
go to step 2

4. If **Finish** [i] == **true** for all i , then the system is in a safe state





Resource-Request Algorithm for Process P_i

$Request_i$ = request vector for process P_i . If **$Request_i[j] = k$** then process P_i wants k instances of resource type R_j


1. If **$Request_i \leq Need_i$** go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If **$Request_i \leq Available$** , go to step 3. Otherwise P_i must wait, since resources are not available
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$Available = Available - Request_i;$

$Allocation_i = Allocation_i + Request_i;$

$Need_i = Need_i - Request_i;$

- If safe \Rightarrow the resources are allocated to P_i
- If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored





Example of Banker's Algorithm

- 5 processes P_1 through P_5 ; $P_2(4,2,1)$
- 3 resource types:
A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 :

<u>Need</u>		<u>Allocation</u>		<u>Max</u>	<u>Available(Work)</u>
A B C		A B C		A B C	A B C
7 4 3	P_1	0 1 0		7 5 3	3 3 2
1 2 2	P_2	2 0 0		3 2 2	
6 0 0	P_3	3 0 2		9 0 2	
0 1 1	P_4	2 1 1		2 2 2	
4 3 1	P_5	0 0 2		4 3 3	






Example (Cont.)

- The content of the matrix ***Need*** is defined to be ***Need = Max – Allocation***

	<u><i>Need</i></u>		
	<i>A</i>	<i>B</i>	<i>C</i>
P_1	7	4	3
P_2	1	2	2
P_3	6	0	0
P_4	0	1	1
P_5	4	3	1

- The system is in a safe state since the sequence $\langle P_2, P_4, P_5, P_1, P_3 \rangle$ satisfies safety criteria



Example of Banker's Algorithm

- 5 processes P_1 through P_5 ;
3 resource types:
A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 :

<u>Need</u>		<u>Allocation</u>	<u>Max</u>	<u>Available(Work)</u>
A B C		A B C	A B C	A B C
7 4 3	P_1	0 1 0	7 5 3	3 3 2
1 2 2	P_2	2 0 0	3 2 2	5 3 2
6 0 0	P_3	3 0 2	9 0 2	7 4 3
0 1 1	P_4	2 1 1	2 2 2	7 4 5
4 3 1	P_5	0 0 2	4 3 3	7 5 5 10 5 7

Available = Total resources - Total allocation
 $= (10, 5, 7) - (7, 2, 5) = (3, 3, 2)$

Safe Sequence = $\langle P_2, P_4, P_5, P_1, P_3 \rangle$




Example: P_1 Request (1,0,2)

- Check that Request \leq Available (that is, $(1,0,2) \leq (3,3,2) \Rightarrow$ true


	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement
- Can request for (3,3,0) by P_4 be granted?
- Can request for (0,2,0) by P_0 be granted?



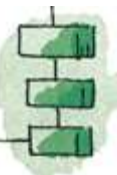
Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme




Single Instance of Each Resource Type


- Maintain **wait-for** graph
 - Nodes are processes
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- An algorithm to detect a cycle in a graph requires an order of n^2 operations, where n is the number of vertices in the graph




Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph




Corresponding wait-for graph

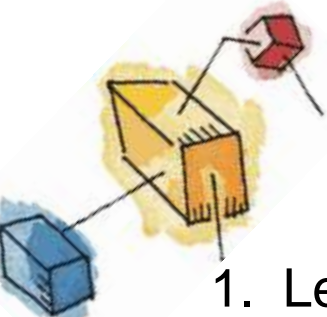


Several Instances of a Resource Type

- **Available:** A vector of length m indicates the number of available resources of each type
- **Allocation:** An $n \times m$ matrix defines the number of resources of each type currently allocated to each process
- **Request:** An $n \times m$ matrix indicates the current request of each process. If **Request** $[i][j] = k$, then process P_i is requesting k more instances of resource type R_j .



Detection Algorithm



1. Let **Work** and **Finish** be vectors of length m and n , respectively
Initialize:

(a) **Work** = **Available**

(b) For $i = 1, 2, \dots, n$, if **Allocation** $_i \neq 0$, then
Finish $[i] = \text{false}$; otherwise, **Finish** $[i] = \text{true}$


2. Find an index i such that both:

(a) **Finish** $[i] == \text{false}$

(b) **Request** $_i \leq \text{Work}$

If no such i exists, go to step 4






Detection Algorithm (Cont.)


3. **$Work = Work + Allocation_i$**
 $Finish[i] = true$
go to step 2

4. If **$Finish[i] == false$** , for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if **$Finish[i] == false$** , then P_i is deadlocked





Recovery from Deadlock: Process Termination

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated
- In which order should we choose to abort?
 1. Priority of the process
 2. How long process has computed, and how much longer to completion
 3. Resources the process has used
 4. Resources process needs to complete
 5. How many processes will need to be terminated



Recovery from Deadlock: Resource Preemption

- **Selecting a victim** – minimize cost
- **Rollback** – return to some safe state, restart process for that state
- **Starvation** – same process may always be picked as victim, include number of rollback in cost factor



Dining Philosophers Problem




Figure 6.11 Dining Arrangement for Philosophers

Dining Philosophers Problem



```
semaphore s[5]=1;
```

```
void philosopher (void)
```

```
{
```

```
while(true)
```

```
{
```

```
Thinking();
```

```
wait(take_fork(si));
```

```
wait(take_fork((si+1)%N));
```

```
Eat();    // Critical Section
```

```
signal(put_fork(si));
```

```
signal(put_fork((si+1)%N));
```

```
}
```

```
}
```

S0	S1	S2	S3	S4
1	1	1	1	1

P0 – f0 , f1 – S0 , S1


P1 - f1 , f2 - S1 , S2

P2 - f2 , f3 - S2 , S3

P3 – f3 , f4 - S3 , S4

P4 - f4 , f0 - S4, S0





Consider the following snapshot of a system

Answer the following using Banker's Algorithm

- What is the content of Need Matrix?
- Is the system in safe state?
- If the request from P1 arrives for (0, 4, 2, 0); can the request be granted immediately?

	Allocation				Max				Available			
	A	B	C	D	A	B	C	D	A	B	C	D
P0	0	0	1	2	0	0	1	2	1	5	2	0
P1	1	0	0	0	1	7	5	0				
P2	1	3	5	4	2	3	5	6				
P3	0	6	3	2	0	6	5	2				
P4	0	0	1	4	0	6	5	6				

